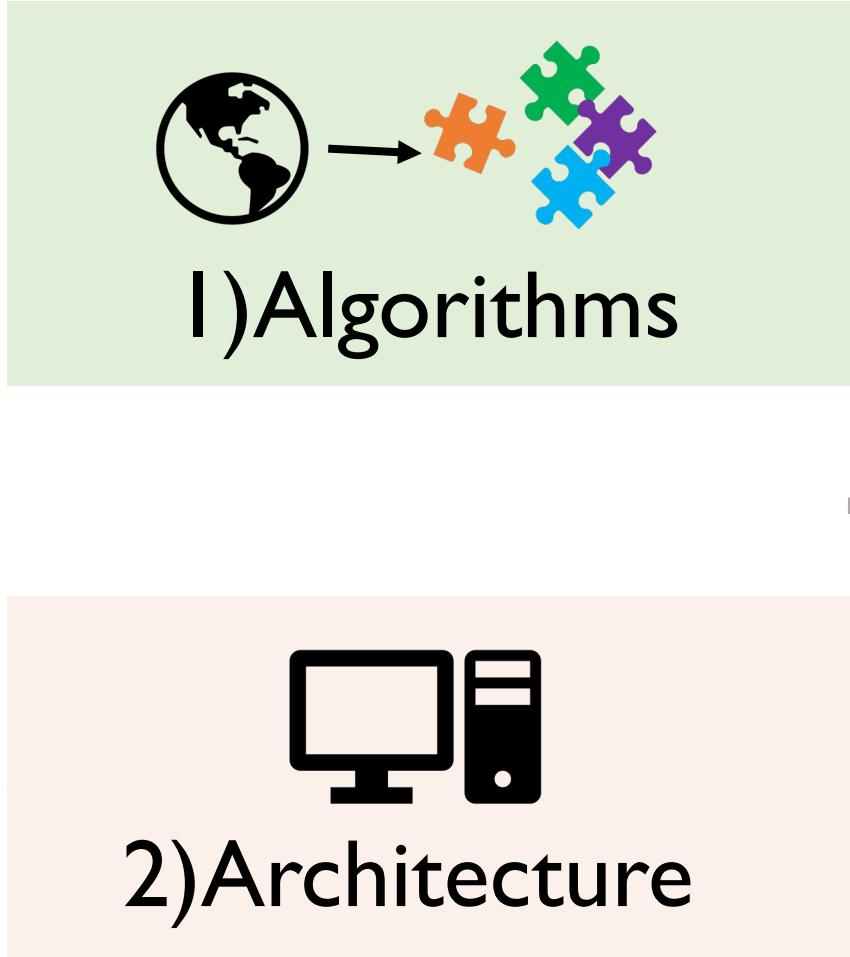


HPC



Solutions

3) Shared Memory
OpenMP

4) Distributed Memory
MPI

5) Unified Engine
Spark

What was challenging about the previous approaches?

3)Unified Engine

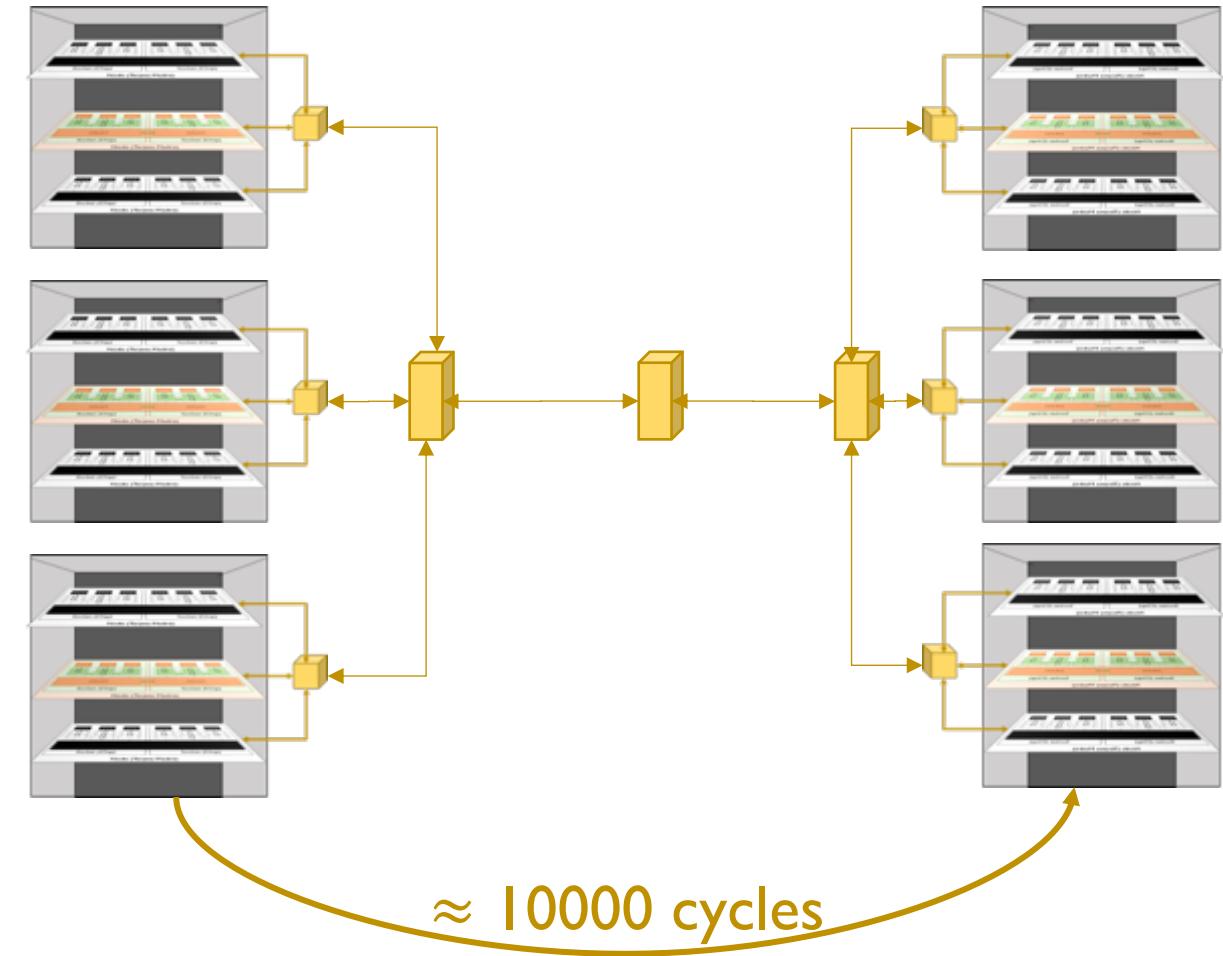
Spark + Hadoop



CARTOONSTOCK

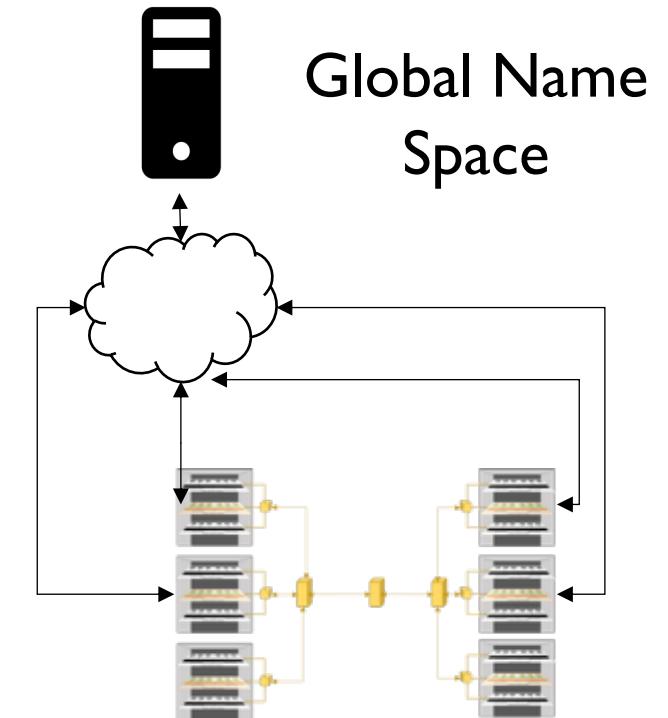
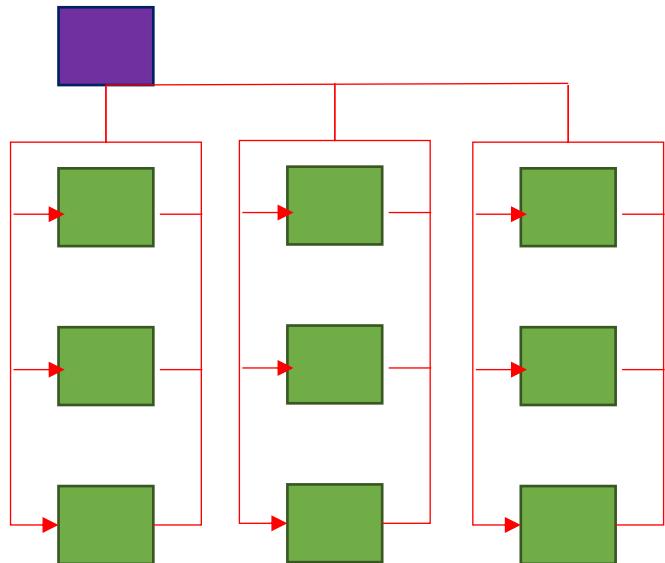
Search ID: aton3981

Distributed Memory



Spark

Fine granularity



*Parallel Regions in a distributed
memory ecosystem*

What do we need in a unified engine?

Distribute instructions

Unified Data Management

Instructions MapReduce / Spark

Fine Level of granularity

I. Transformations & Actions

2. Resilient Distributed Datasets

Data Management HDFS

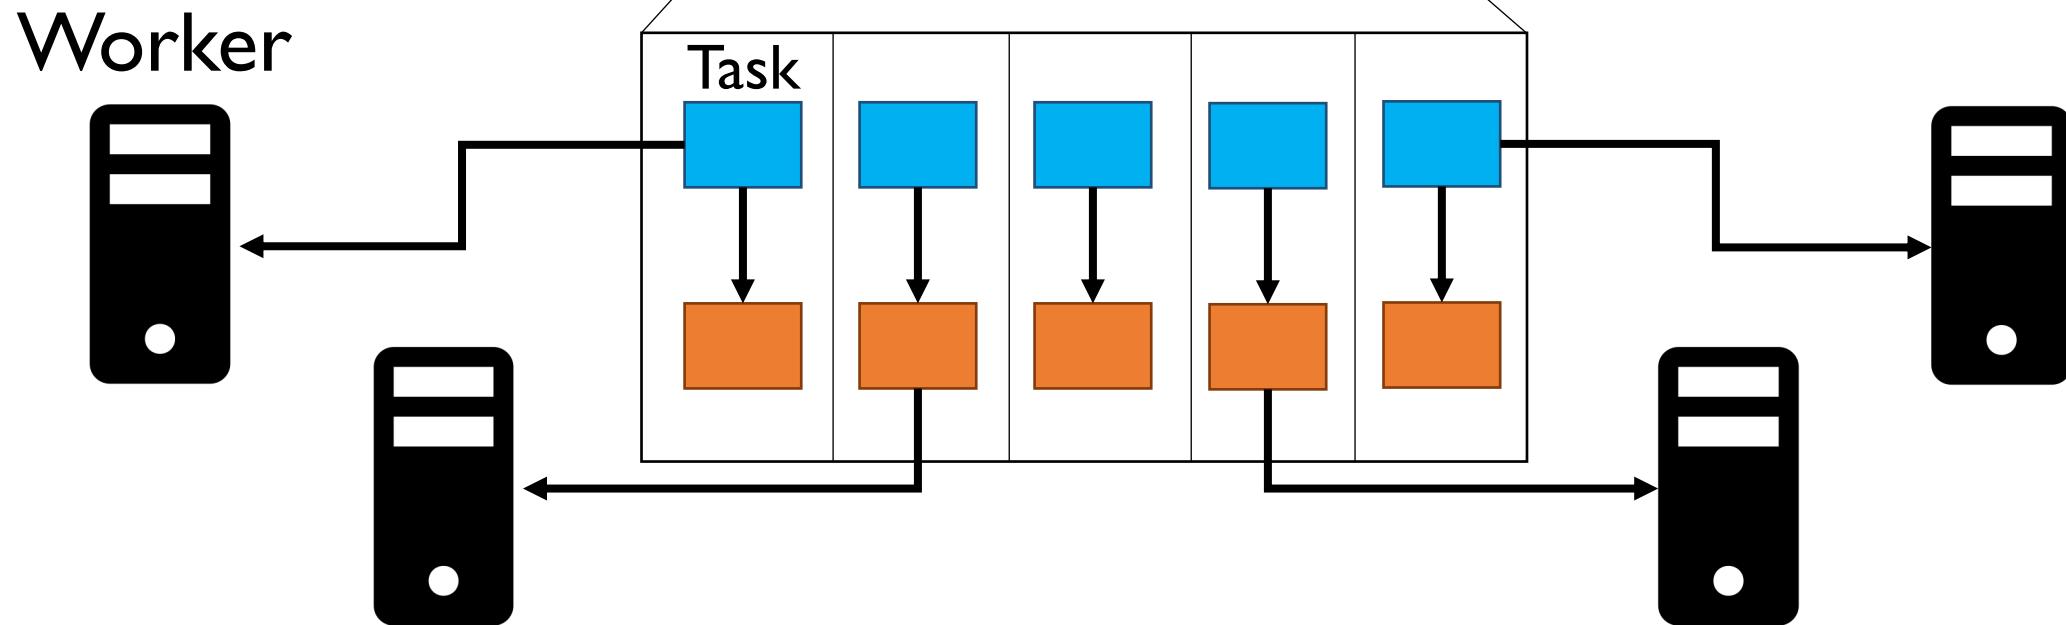
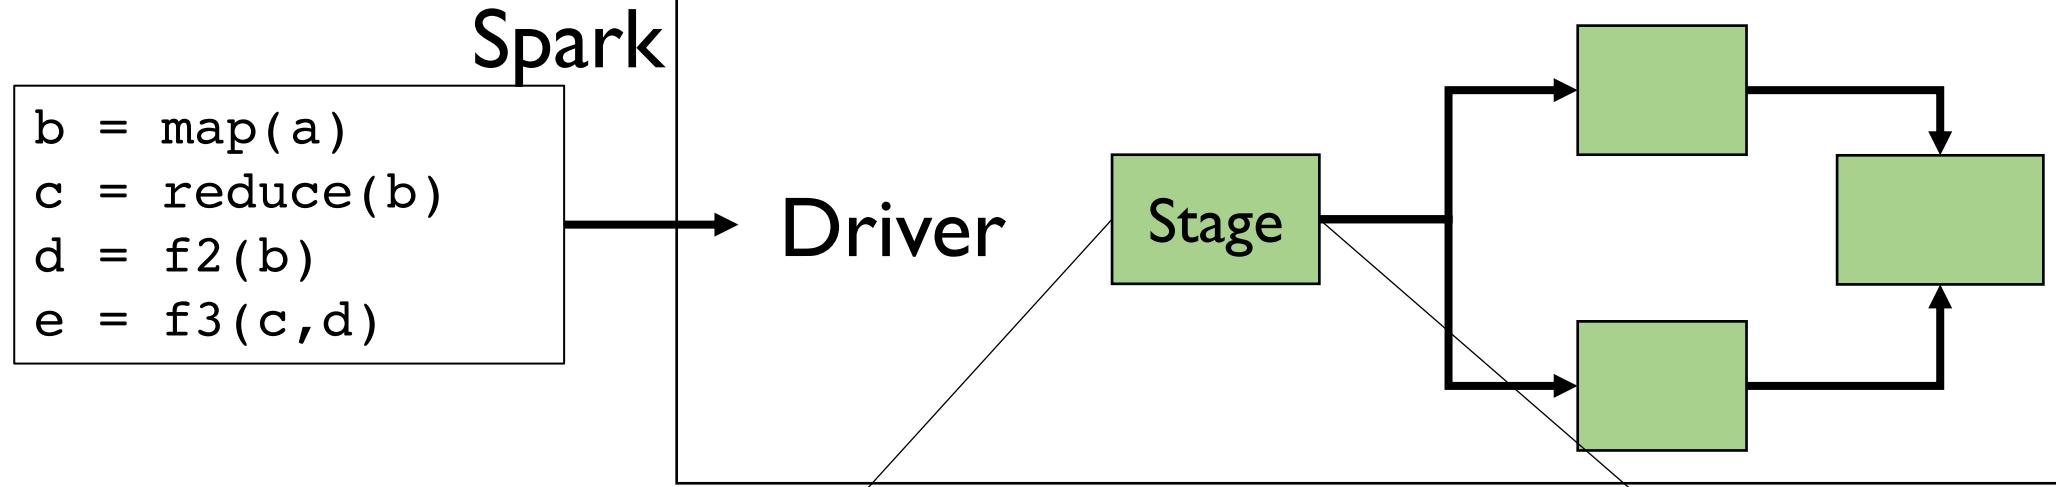
Global naming Space

Fault Tolerance & Portability

Cluster Manager: YARN

Resource Allocation & Job Scheduling

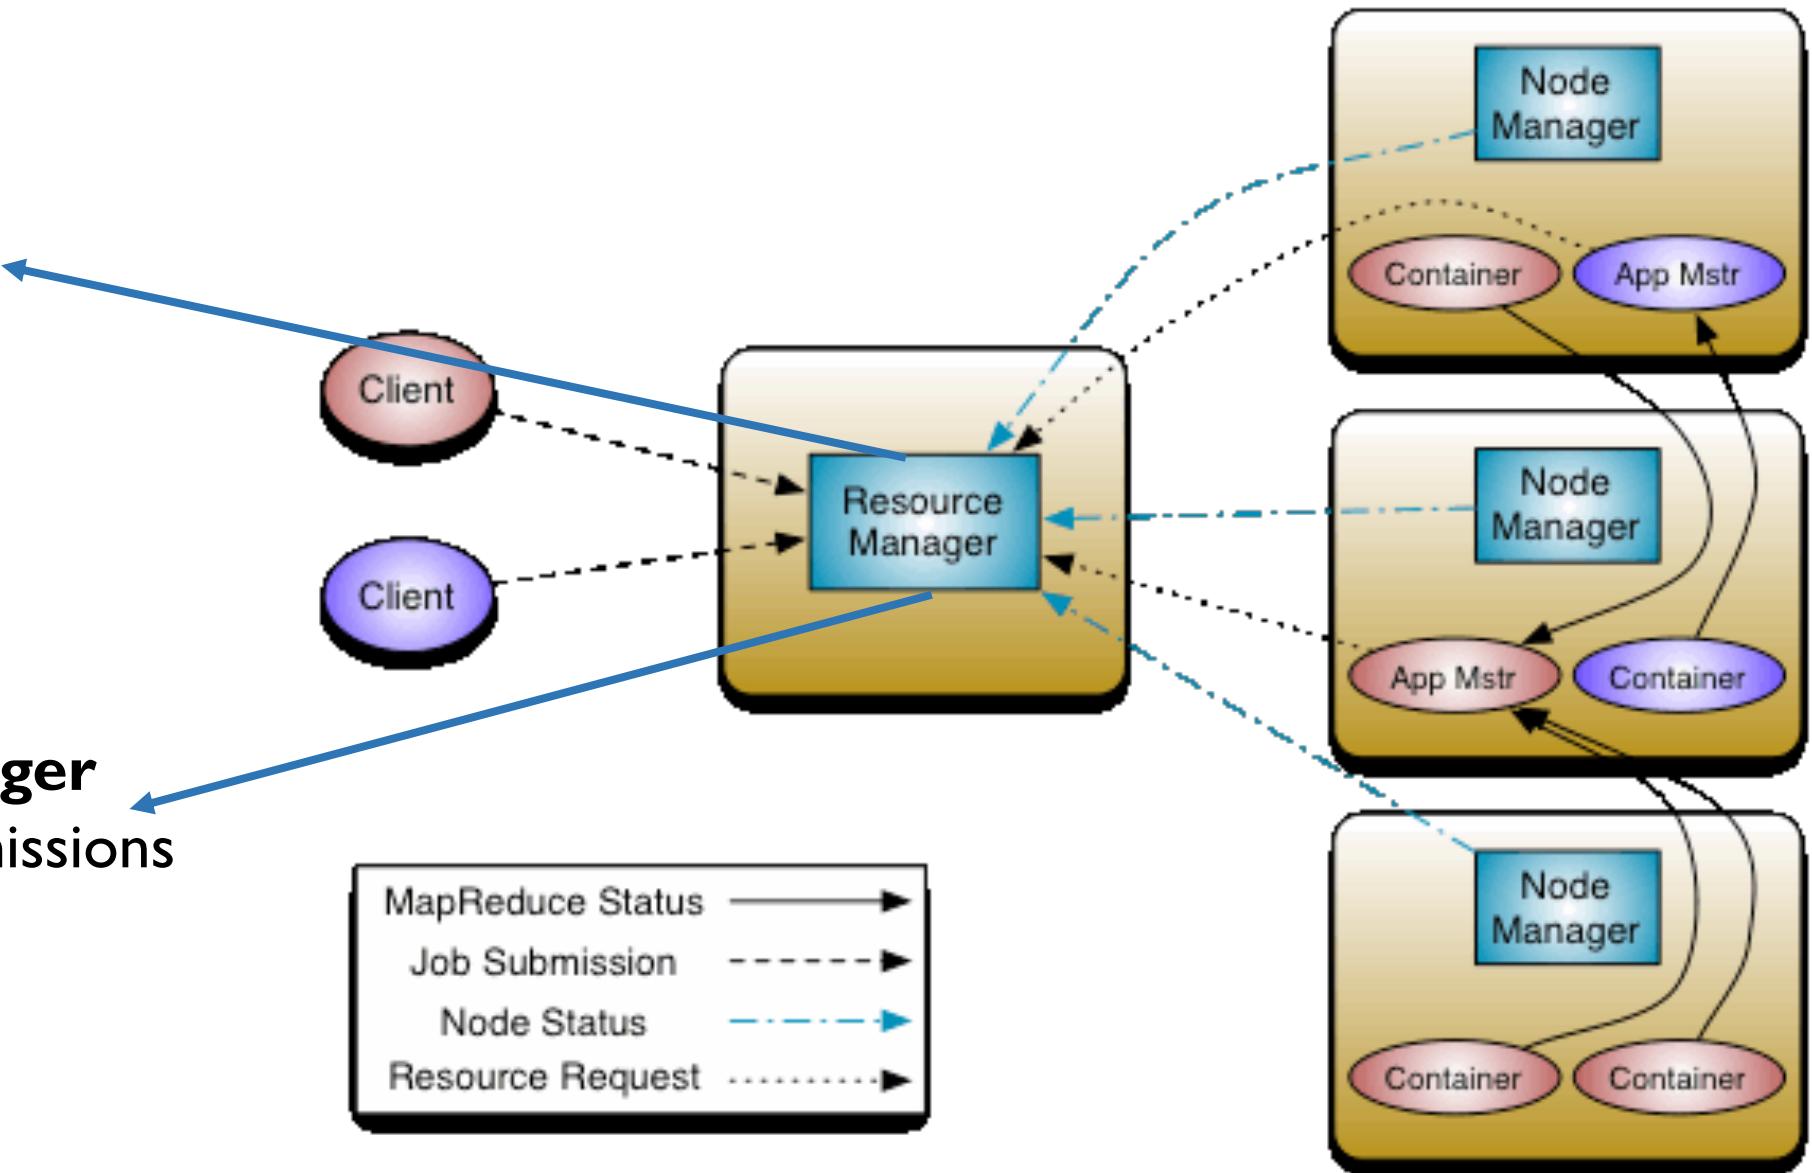
Spark in a nutshell



YARN: Provides the executors

Scheduler
Allocates Resources

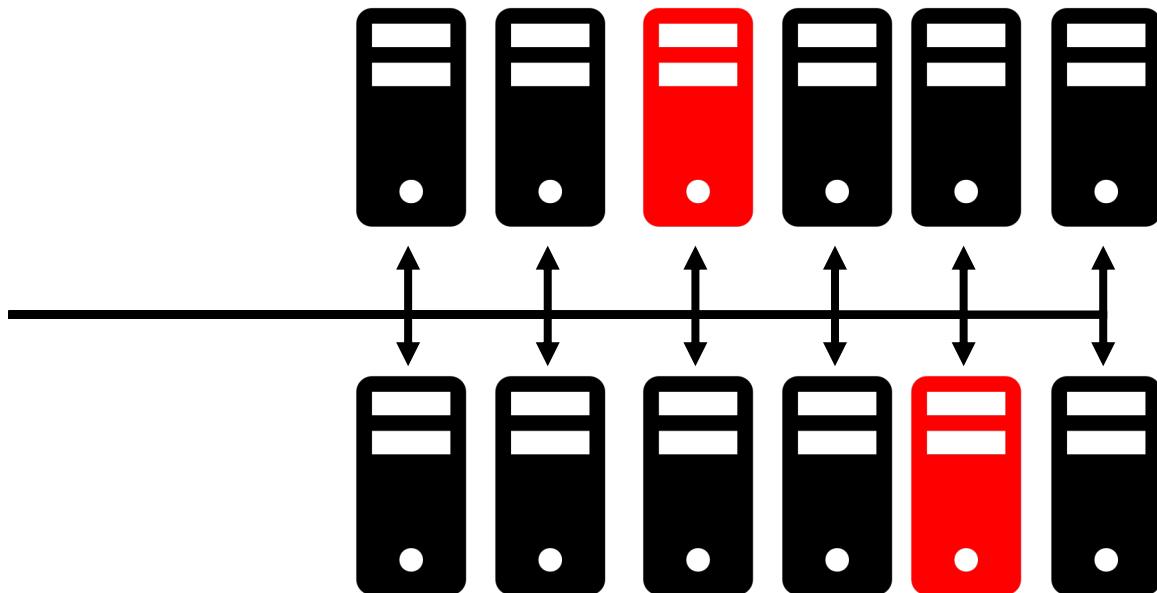
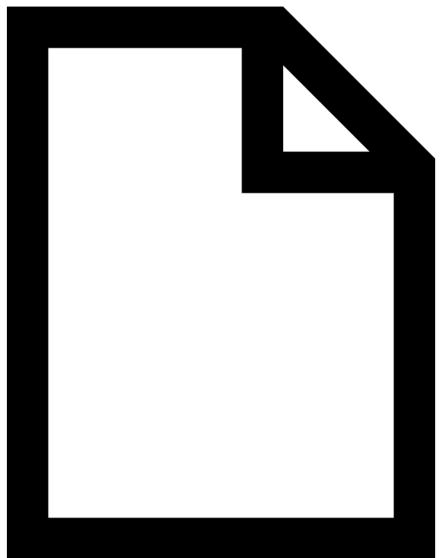
ApplicationsManager
Negotiates job-submissions
Recovery & Failure



<https://hadoop.apache.org/docs/current/hadoop-yarn/hadoop-yarn-site/YARN.html>

What is Hadoop?

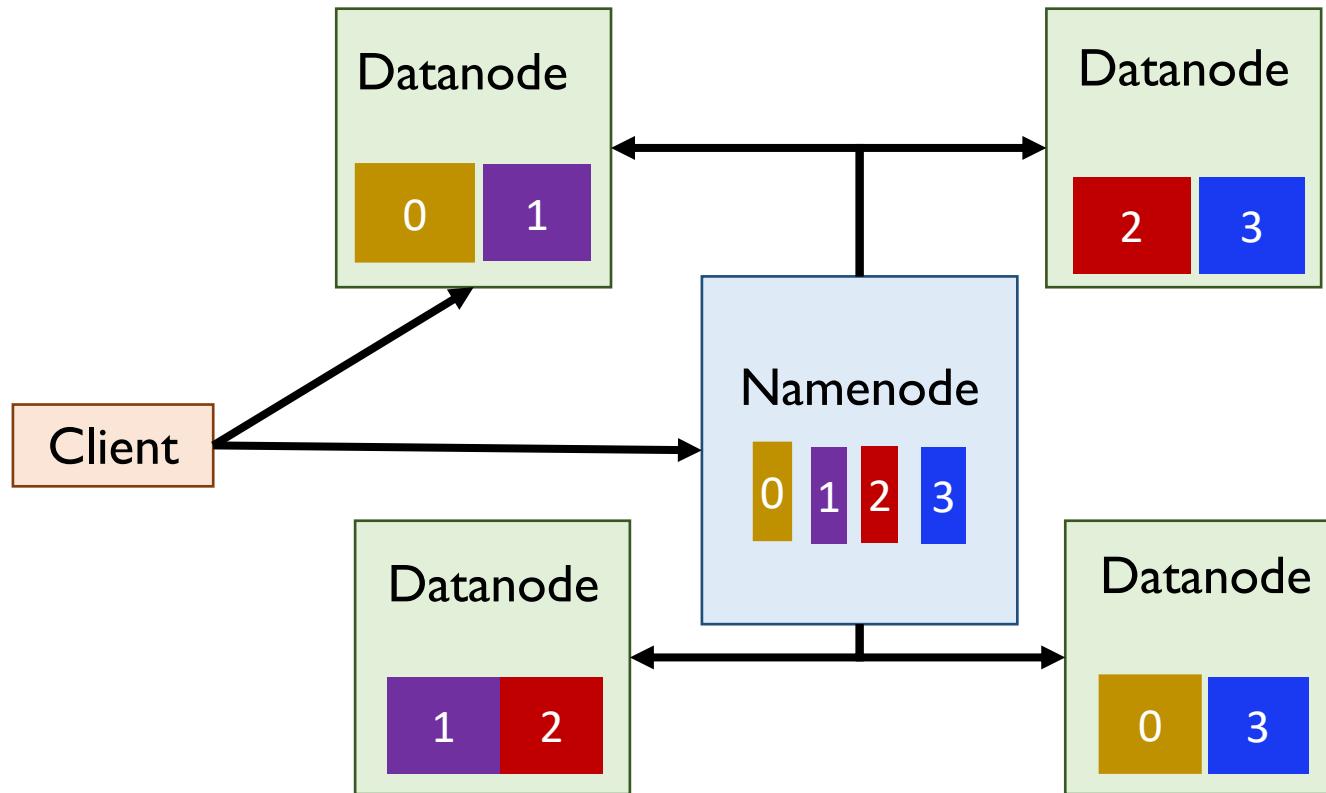
Library in Java to manage large data sets in clusters with failures



<https://hadoop.apache.org>

<https://spark.apache.org>

HDFS: Hadoop File System

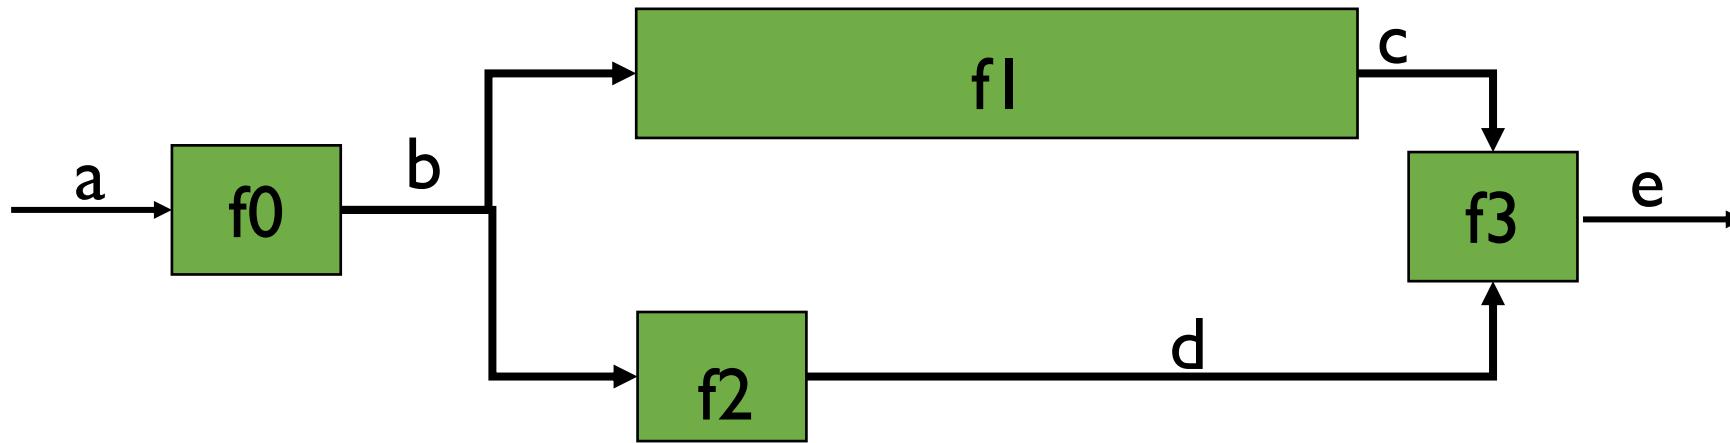


Fault Tolerance

Moving Computation less expensive than moving data

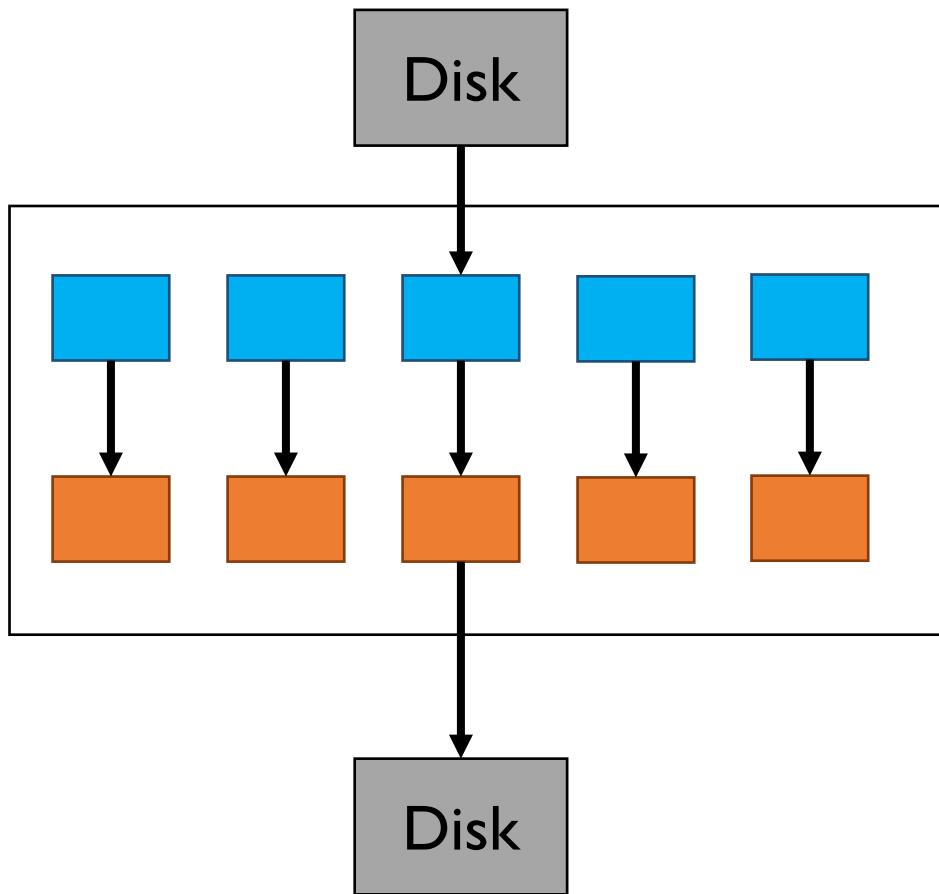
Portability

Why Spark is successful: Functional Programming

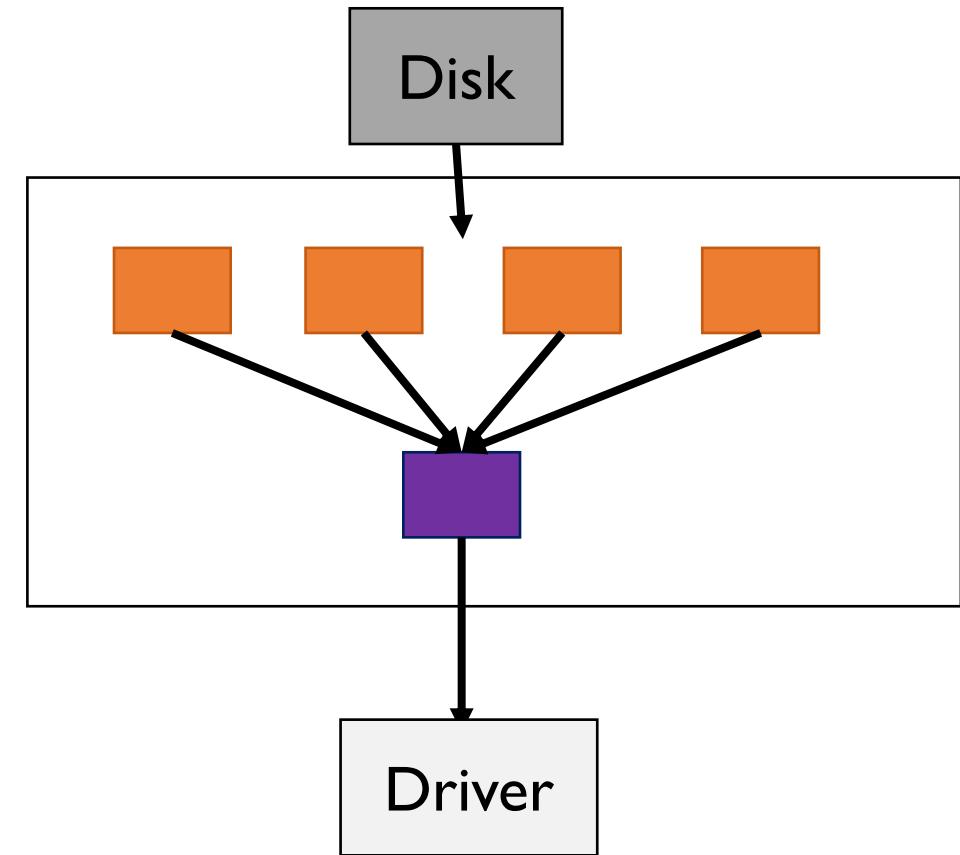


Map Reduce Paradigm

map



reduce

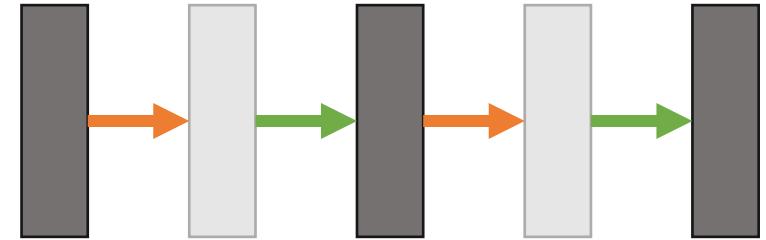


Spark: Resilient Distributed Datasets (RDDs)

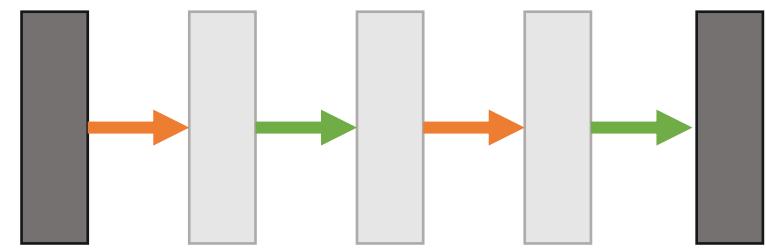
```
for i in 1:n  
  b = map(a)  
  a = reduce(b)
```



Map Reduce



Spark



How to create RDDS

`parallelize(data)`

Distribute data in driver program

`textFile("data.txt")...`

Distribute dataset in external storage: HDFS, shared filesystem ...

Accessible to all workers

How to **persist** in memory

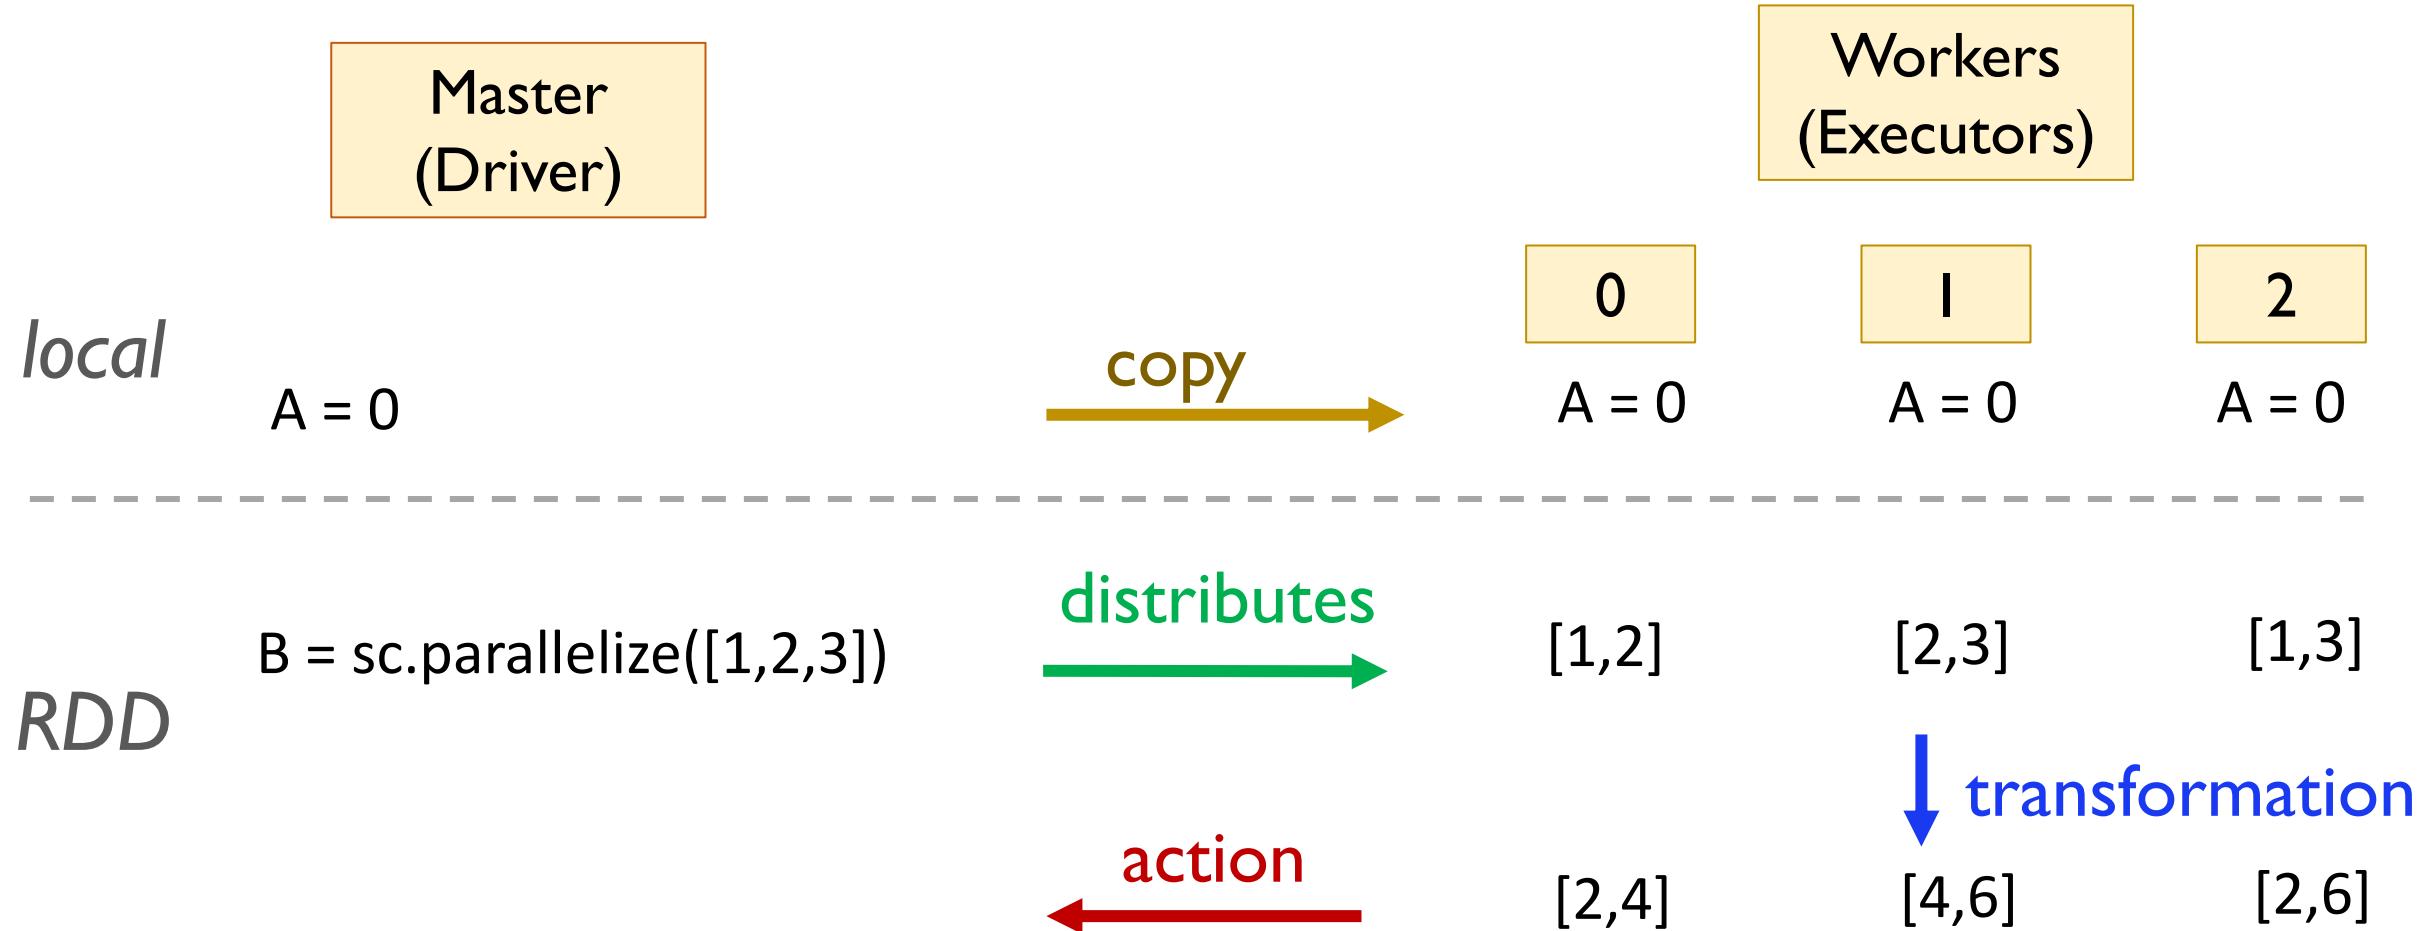
Default (*Lazy Operation*)

Compute transformation on RDDs each time an action is required

`persist(Storage_level)...`

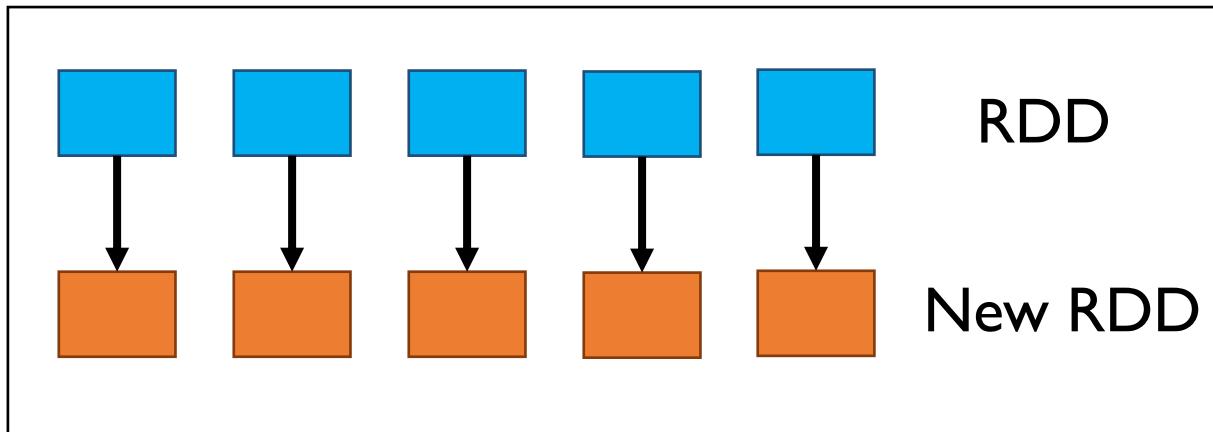
Keeps in memory RDD without recomputing each time

Local vs Global Variables



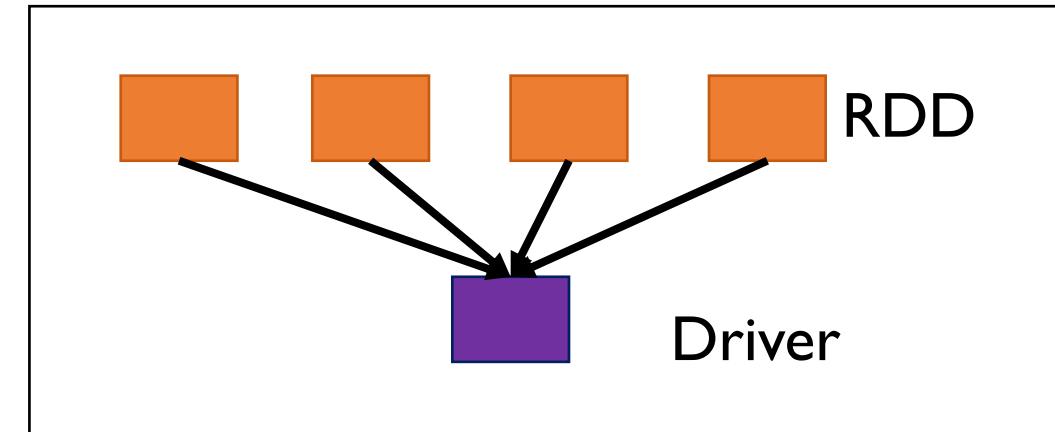
Spark Enhancing

Transformation



`map(func)`
`filter(func)`
`flatMap(func)`
`distinct()`
`sample(k)`
`foreach(func)`

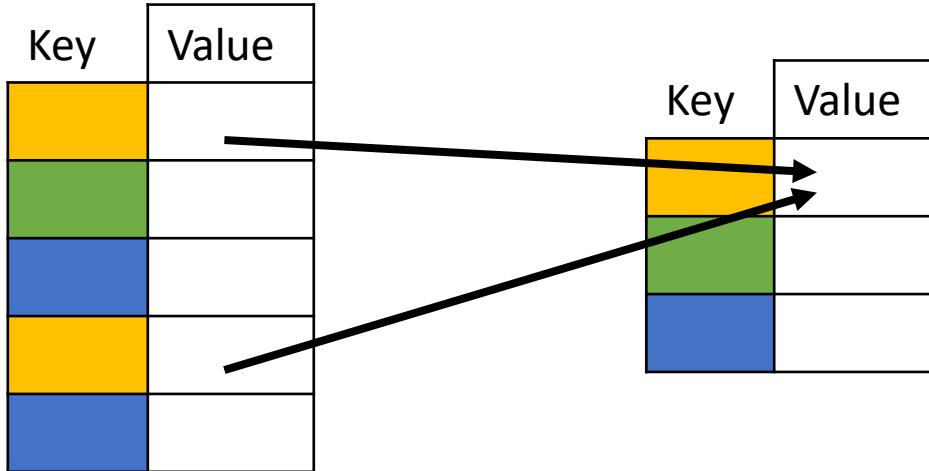
Action



`collect()`
`take(k)`
`count()`
`reduce()`
`takeSample(k)`
`saveAsTextFile(...)`

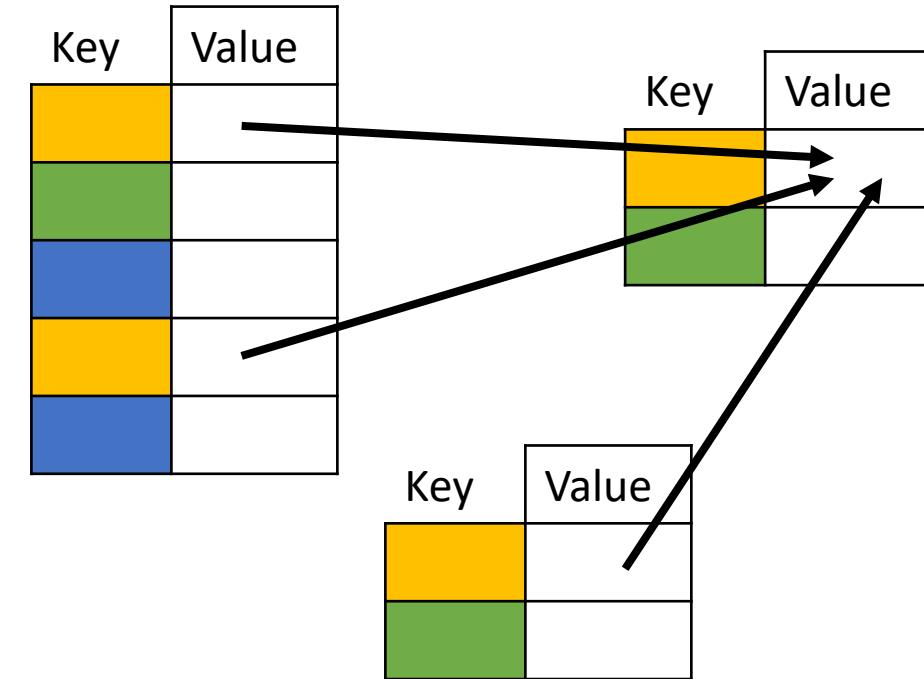
Spark RDD operations

Key-value operations



`reduceByKey(func)`
`groupByKey()`
`sortByKey()`

Interaction with other RDD



`join(other)`
`cogroup(other)`

Initialize Spark Application

SparkConf()

`setAppName(appName)` Define application name in cluster

`setMaster(master)` URL of cluster manager

SparkContext(conf) Access to cluster (No needed in the terminal)

Cluster Options for Spark

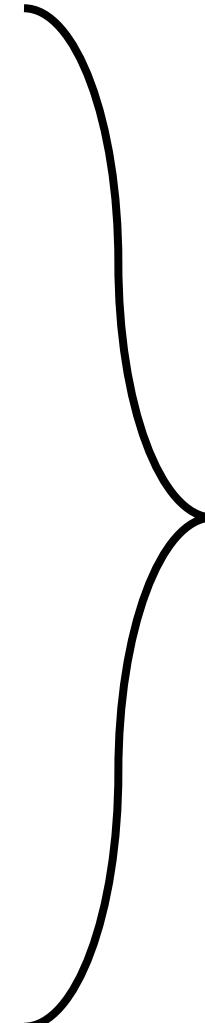
Local

Standalone

Apache Mesos

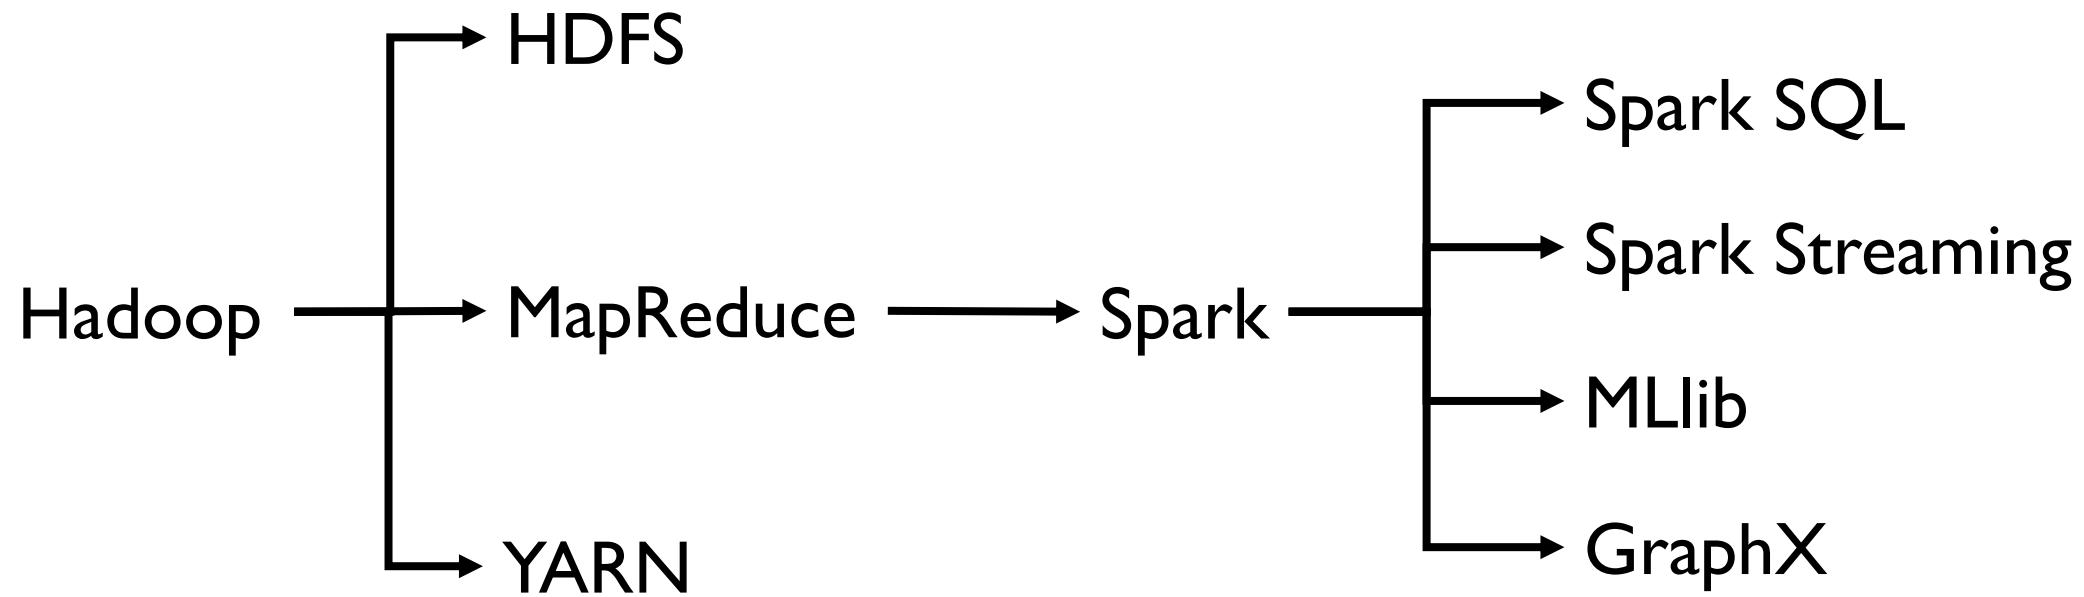
Hadoop YARN

Kubernetes



Scheduling

Monitoring



<https://hadoop.apache.org>

<https://spark.apache.org>

How to run my application?

- 1) Have a Apache Spark distribution, usually prebuilt for Apache Hadoop
- 2) Choose programming language:

Java, Scala, **Python**, R, SQL

- 3) In file add library :

```
from pyspark.sql import SparkSession
```

- 4) Execute defining number of cores involve (to run locally)

```
bin/spark-submit --master local[4] hello_world_spark.py
```

Monitoring

- Create folder `/tmp/spark-events` (only once)
- Execute with flag

```
bin/spark-submit --master local[4] --conf spark.eventLog.enabled=true pi.py
```

- Recover post-mortem data

```
./sbin/start-history-server.sh
```

- Check the info at <http://localhost:18080>
- During execution you can check <http://localhost:4040>

PythonPageRank - Details for Job 0 - Mozilla Firefox

PythonPageRank - Details × +

localhost:18080/history/local-1534345419942/jobs/job/?id=0

Search

Apache Spark 2.3.1 Jobs Stages Storage Environment Executors PythonPageRank application UI

Details for Job 0

Status: SUCCEEDED
Completed Stages: 23

- Event Timeline
- DAG Visualization

Completed Stages (23)

| Stage Id | Description | Submitted | Duration | Tasks: Succeeded/Total | Input | Output | Shuffle Read | Shuffle Write |
|----------|---|------------------------------|----------|------------------------|---------|--------|--------------|---------------|
| 22 | collect at /home/user/spark-2.3.1-bin-hadoop2.7/examples/src/main/python/pagerank.py:85 | +details 2018/08/15 12:03:49 | 0.3 s | 11/11 | | | 432.0 B | |
| 21 | reduceByKey at /home/user/spark-2.3.1-bin-hadoop2.7/examples/src/main/python/pagerank.py:82 | +details 2018/08/15 12:03:49 | 0.3 s | 11/11 | | | 962.0 B | 432.0 B |
| 20 | join at /home/user/spark-2.3.1-bin-hadoop2.7/examples/src/main/python/pagerank.py:78 | +details 2018/08/15 12:03:49 | 0.3 s | 11/11 | 235.0 B | | 248.0 B | 962.0 B |
| 19 | reduceByKey at /home/user/spark-2.3.1-bin-hadoop2.7/examples/src/main/python/pagerank.py:82 | +details 2018/08/15 12:03:48 | 0.3 s | 10/10 | | | 620.0 B | 248.0 B |
| 18 | join at /home/user/spark-2.3.1-bin-hadoop2.7/examples/src/main/python/pagerank.py:78 | +details 2018/08/15 12:03:48 | 0.3 s | 10/10 | 235.0 B | | 248.0 B | 620.0 B |
| 17 | reduceByKey at /home/user/spark-2.3.1-bin-hadoop2.7/examples/src/main/python/pagerank.py:82 | +details 2018/08/15 12:03:48 | 0.2 s | 9/9 | | | 736.0 B | 248.0 B |
| 16 | join at /home/user/spark-2.3.1-bin-hadoop2.7/examples/src/main/python/pagerank.py:78 | +details 2018/08/15 12:03:48 | 0.3 s | 9/9 | 235.0 B | | 432.0 B | 736.0 B |
| 15 | reduceByKey at /home/user/spark-2.3.1-bin-hadoop2.7/examples/src/main/python/pagerank.py:82 | +details 2018/08/15 12:03:47 | 0.2 s | 8/8 | | | 962.0 B | 432.0 B |
| 14 | join at /home/user/spark-2.3.1-bin-hadoop2.7/examples/src/main/python/pagerank.py:78 | +details 2018/08/15 12:03:47 | 0.3 s | 8/8 | 235.0 B | | 301.0 B | 962.0 B |
| 13 | reduceByKey at /home/user/spark-2.3.1-bin-hadoop2.7/examples/src/main/python/pagerank.py:82 | +details 2018/08/15 12:03:47 | 0.2 s | 7/7 | | | 848.0 B | 301.0 B |
| 12 | join at /home/user/spark-2.3.1-bin-hadoop2.7/examples/src/main/python/pagerank.py:78 | +details 2018/08/15 12:03:47 | 0.2 s | 7/7 | 235.0 B | | 248.0 B | 848.0 B |
| 11 | reduceByKey at /home/user/spark-2.3.1-bin-hadoop2.7/examples/src/main/python/pagerank.py:82 | +details 2018/08/15 12:03:46 | 0.2 s | 6/6 | | | 620.0 B | 248.0 B |
| 10 | join at /home/user/spark-2.3.1-bin-hadoop2.7/examples/src/main/python/pagerank.py:78 | +details 2018/08/15 12:03:46 | 0.3 s | 6/6 | 235.0 B | | 248.0 B | 620.0 B |
| 9 | reduceByKey at /home/user/spark-2.3.1-bin-hadoop2.7/examples/src/main/python/pagerank.py:82 | +details 2018/08/15 12:03:46 | 0.2 s | 5/5 | | | 736.0 B | 248.0 B |
| 8 | join at /home/user/spark-2.3.1-bin-hadoop2.7/examples/src/main/python/pagerank.py:78 | +details 2018/08/15 12:03:46 | 0.2 s | 5/5 | 235.0 B | | 432.0 B | 736.0 B |
| 7 | reduceByKey at /home/user/spark-2.3.1-bin-hadoop2.7/examples/src/main/python/pagerank.py:82 | +details 2018/08/15 12:03:45 | 0.1 s | 4/4 | | | 962.0 B | 432.0 B |
| 6 | join at /home/user/spark-2.3.1-bin-hadoop2.7/examples/src/main/python/pagerank.py:78 | +details 2018/08/15 12:03:45 | 0.2 s | 4/4 | 235.0 B | | 246.0 B | 962.0 B |
| 5 | reduceByKey at /home/user/spark-2.3.1-bin-hadoop2.7/examples/src/main/python/pagerank.py:82 | +details 2018/08/15 12:03:45 | 0.1 s | 3/3 | | | 616.0 B | 246.0 B |
| 4 | join at /home/user/spark-2.3.1-bin-hadoop2.7/examples/src/main/python/pagerank.py:78 | +details 2018/08/15 12:03:45 | 0.1 s | 3/3 | 235.0 B | | 245.0 B | 616.0 B |

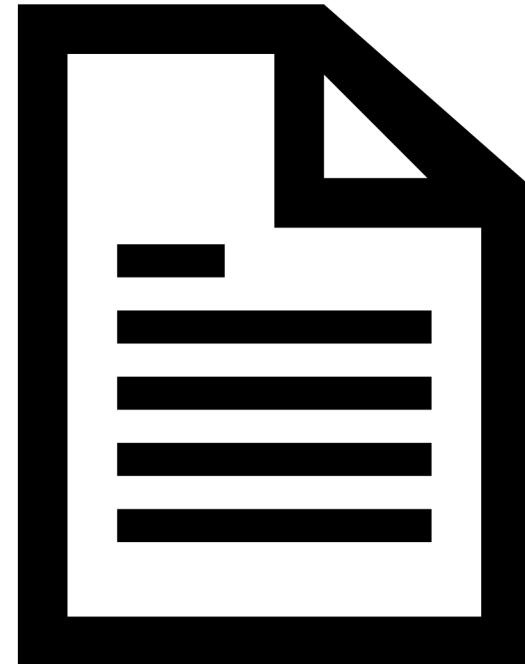
References

<https://spark.apache.org/docs/latest/index.html>

Intro to RDDs, extensions, APIs and deploying

<http://hadoop.apache.org/docs/current/index.html>

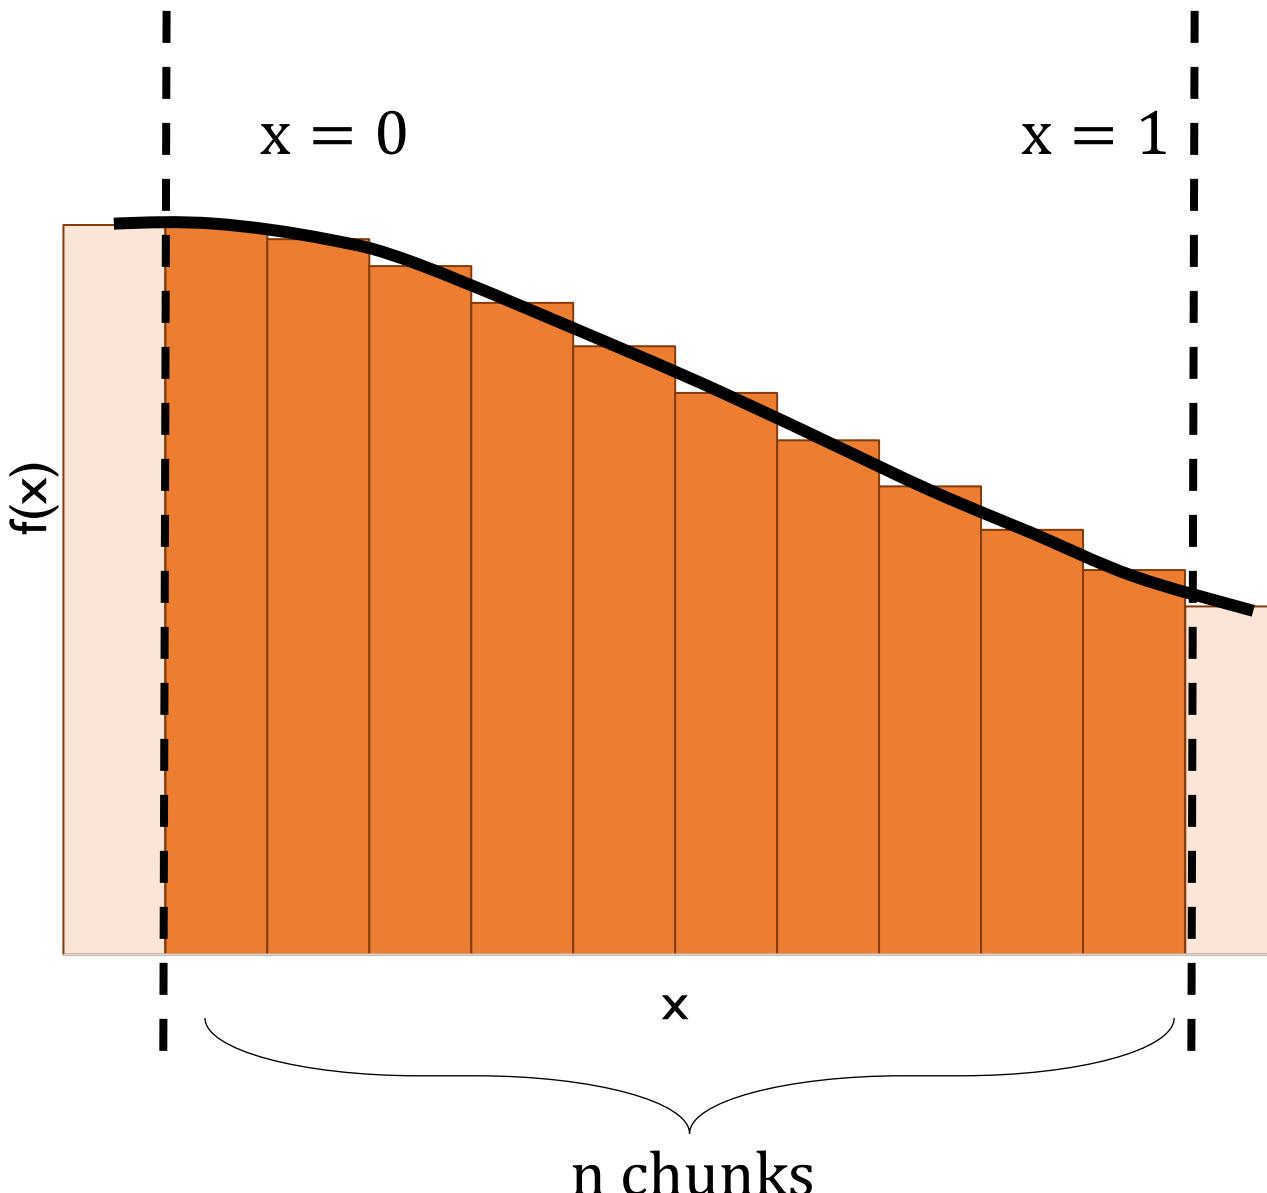
Hadoop overview, including HDFS and YARN



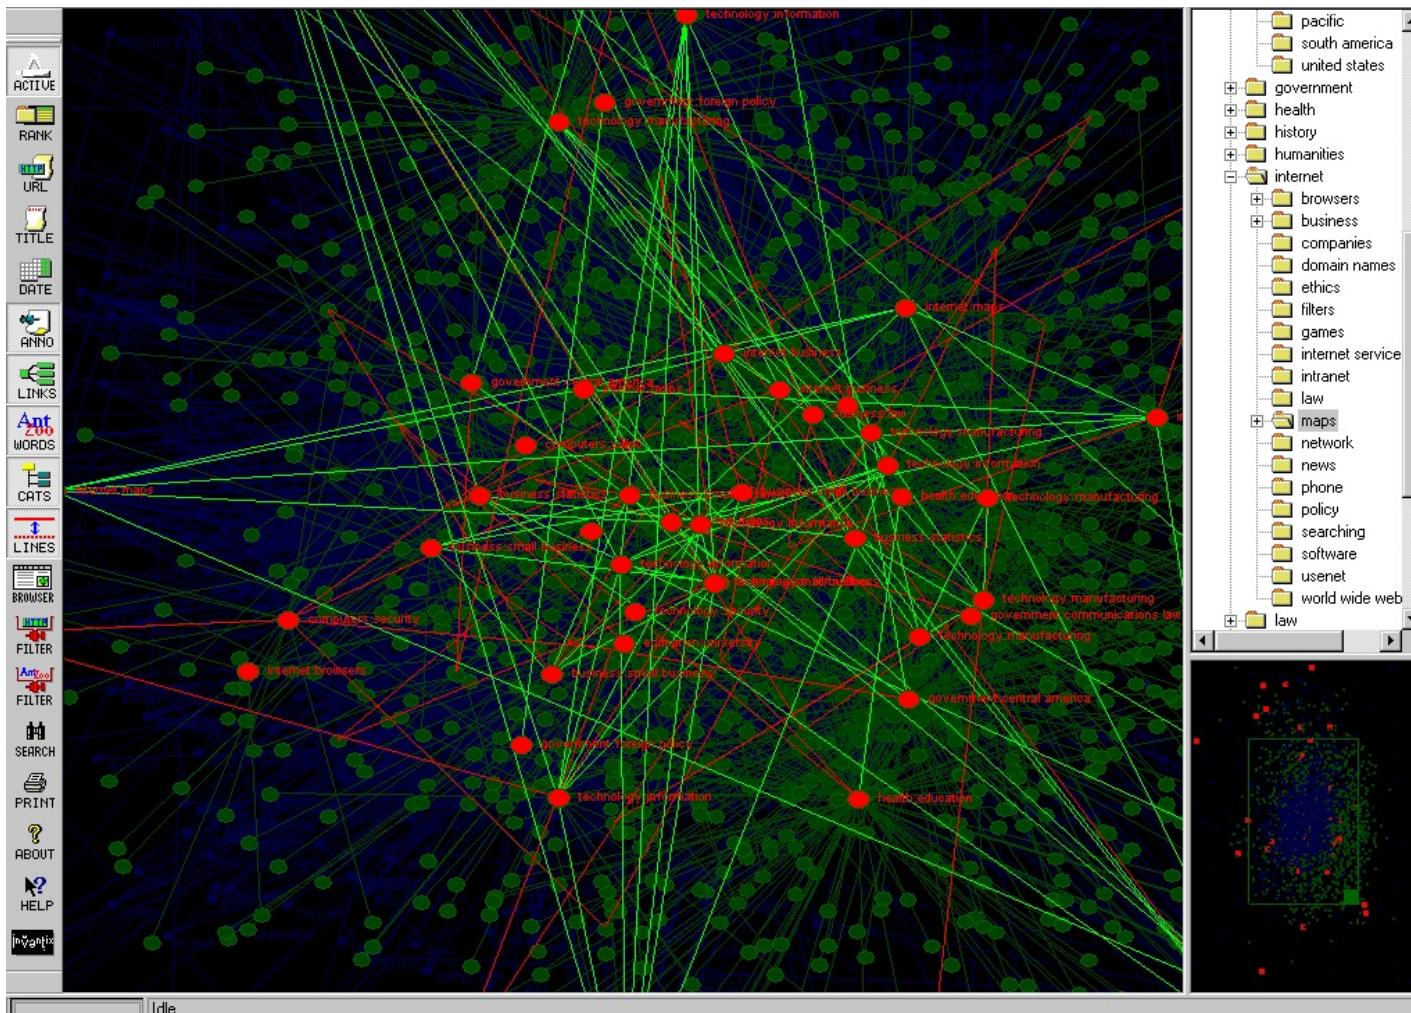
Spark: The Definitive Guide: Big Data Processing
Made Simple, Matei Zaharia, 2018.

Example: Compute Pi ($\pi = 3.14159 \dots$)

$$\pi = \int_0^1 \frac{4}{1+x^2} dx$$

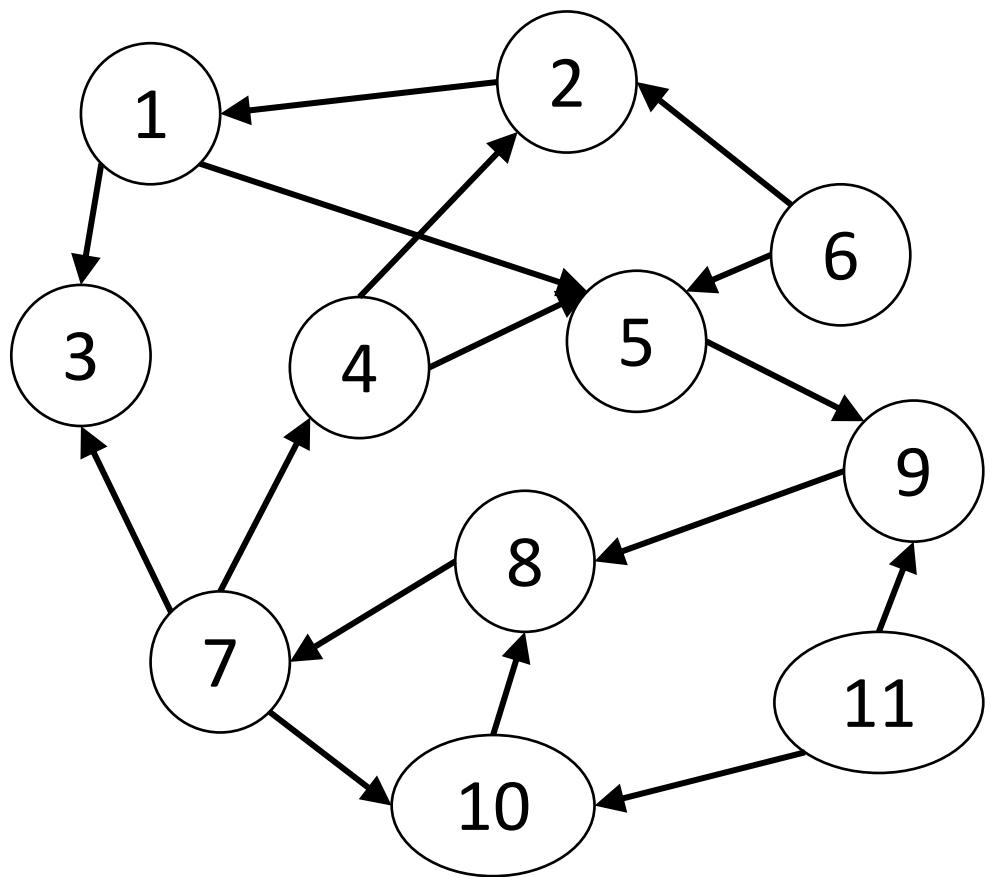


Example: PageRank



An Atlas Of Cyberspaces: Internet Cartographer by Inventix Sotware, mid 2000s

Example: PageRank



$S =$

| | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 |
|----|---|---|---|---|---|---|---|---|---|----|----|
| 1 | | | | | | | | | | | |
| 2 | | | | | | | | | | | |
| 3 | | | | | | | | | | | |
| 4 | | | | | | | | | | | |
| 5 | | | | | | | | | | | |
| 6 | | | | | | | | | | | |
| 7 | | | | | | | | | | | |
| 8 | | | | | | | | | | | |
| 9 | | | | | | | | | | | |
| 10 | | | | | | | | | | | |
| 11 | | | | | | | | | | | |

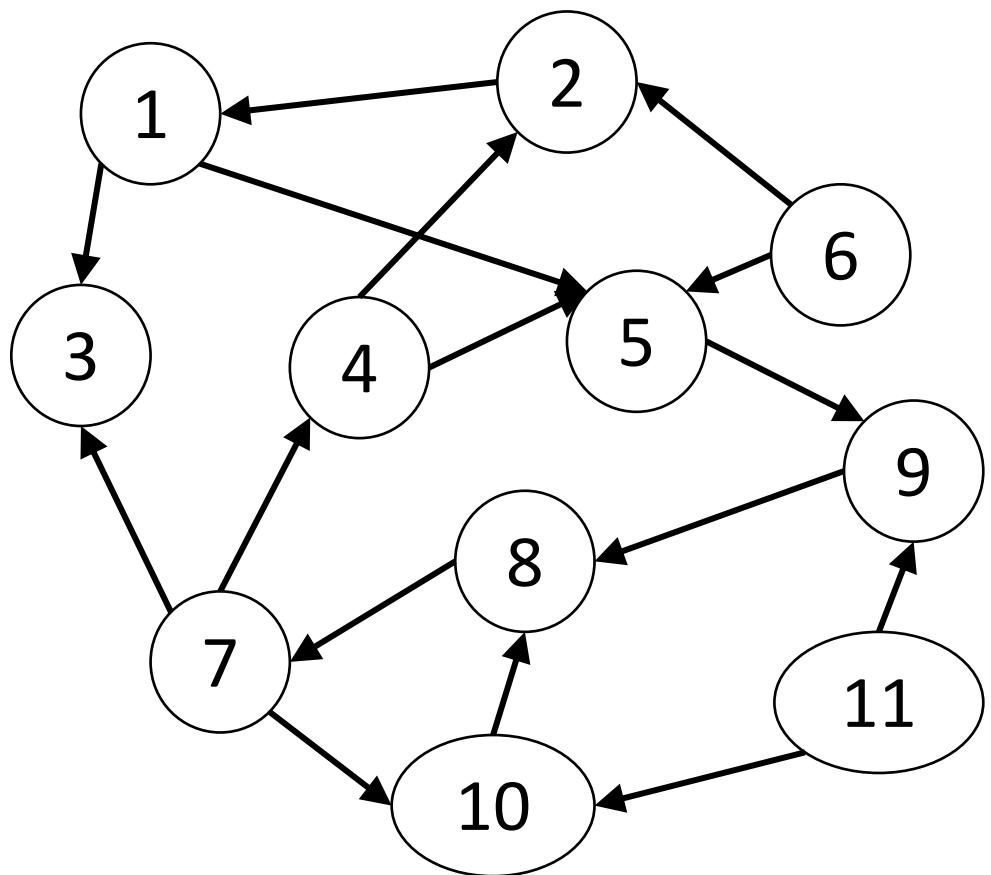
Example: PageRank

Find left eigenvalue $0.85 \cdot S + 0.15 (1 \cdot 1^\top)$

$$r^{(k+1)} = 0.85 \underbrace{S^\top r^{(k)}}_{\text{green bracket}} + 0.15 \cdot 1$$

$$c_j = \sum_{i \text{ ant.of } j} \frac{r_i}{\#\text{outlinks } i}$$

Example: PageRank $(S^T r)_9$?

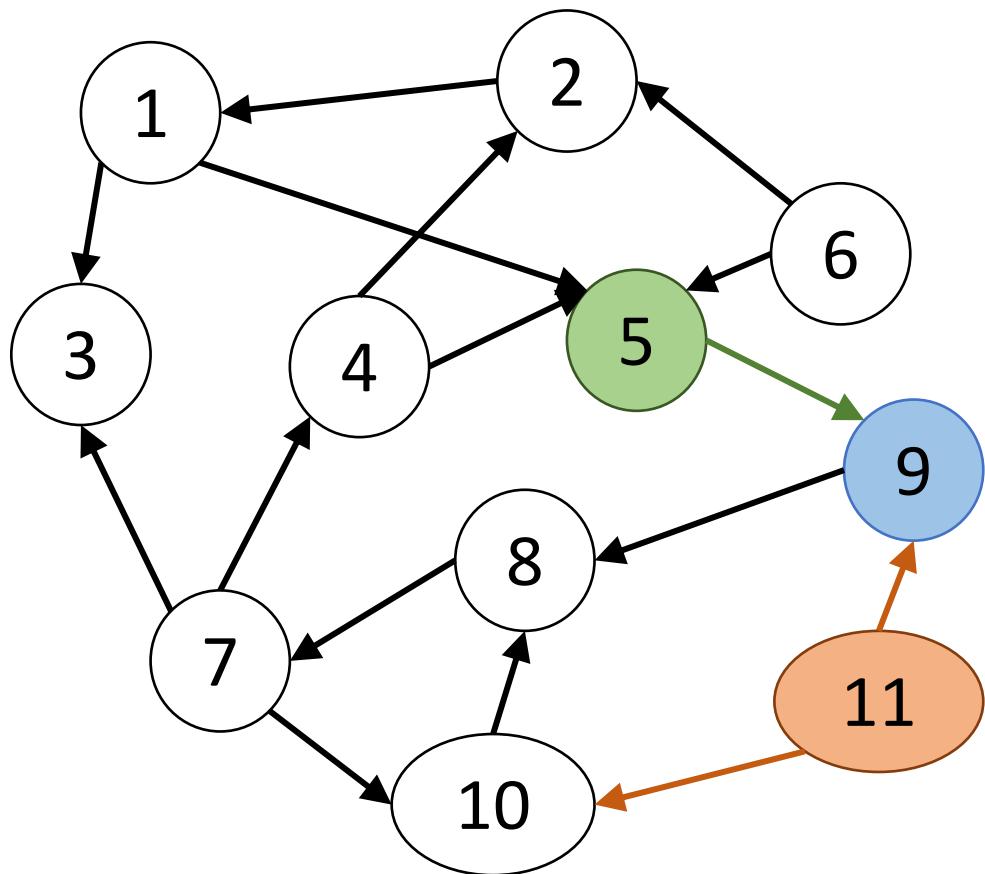


$S =$

| | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 |
|----|---|---|---|---|---|---|---|---|---|----|----|
| 1 | | | | | | | | | | | |
| 2 | | | | | | | | | | | |
| 3 | | | | | | | | | | | |
| 4 | | | | | | | | | | | |
| 5 | | | | | | | | | | | |
| 6 | | | | | | | | | | | |
| 7 | | | | | | | | | | | |
| 8 | | | | | | | | | | | |
| 9 | | | | | | | | | | | |
| 10 | | | | | | | | | | | |
| 11 | | | | | | | | | | | |

Example: PageRank

$$(S^T r)_9 ? = \frac{r_5}{1} + \frac{r_{11}}{2}$$



$S =$

| I | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | II |
|----|---|---|---|---|---|---|---|---|---|----|----|
| I | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 2 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 3 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 4 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 5 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 |
| 6 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 |
| 7 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 |
| 8 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 |
| 9 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 |
| 10 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 |
| II | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 |

| antecedent | descendant |
|------------|------------|
| 1 | [3,5] |
| 2 | [1] |
| 3 | [] |
| 4 | [2,5] |
| 5 | [9] |
| 6 | [2,5] |
| 7 | [3,4,10] |
| 8 | [7] |
| 9 | [8] |
| 10 | [8] |
| 11 | [9,10] |
| Node | r_i |
| 1 | 1 |
| 2 | 1 |
| 3 | 1 |
| 4 | 1 |
| 5 | 1 |
| 6 | 1 |
| 7 | 1 |
| 8 | 1 |
| 9 | 1 |
| 10 | 1 |
| 11 | 1 |

$$i \rightarrow j : c_i(j) = \frac{r_i}{\#\text{outlinks } i}$$

$$u_j = \sum_{i \rightarrow j} c_i(j)$$

$$r_j = 0.85u_j + 0.15$$

| antecedent | contribution |
|------------|------------------------------|
| 1 | [[3,1/2], [5,1/2]] |
| 2 | [[1,1]] |
| 3 | [] |
| 4 | [[2,1/2], [5,1/2]] |
| 5 | [[9,1]] |
| 6 | [[2,1/2], [5,1/2]] |
| 7 | [[3,1/3], [4,1/3], [10,1/3]] |
| 8 | [[7,1]] |
| 9 | [[8,1]] |
| 10 | [[8,1]] |
| 11 | [[9,1/2], [10,1/2]] |

| descendant | contribution |
|------------|-----------------|
| 1 | 1 |
| 2 | 1/2 + 1/2 |
| 3 | 1/2 + 1/3 |
| 4 | 1/3 |
| 5 | 1/2 + 1/2 + 1/2 |
| 6 | 0 |
| 7 | 1 |
| 8 | 1 + 1 |
| 9 | 1 + 1/2 |
| 10 | 1/2 + 1/3 |
| 11 | 0 |

Reduce by key

| Node | r_i |
|------|-------|
| 1 | 1 |
| 2 | 1 |
| 3 | 1 |
| 4 | 0.43 |
| 5 | 1.425 |
| 6 | 0.15 |
| 7 | 1 |
| 8 | 1.85 |
| 9 | 1.425 |
| 10 | 0.86 |
| 11 | 0 |

map